

IN(1):

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

IN(2):

```
df=pd.read_csv('/content/Churn_Modelling.csv')
```

IN(6):

df

OP(6):

```
Out[6]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.88	1
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.58	0
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	113931.57	1
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	93826.63	0
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	79084.10	0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
9995	9996	15606229	Obijaku	771	France	Male	39	5	0.00	2	1	0	96270.64	0
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61	1	1	1	101699.77	0
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	1	0	1	42085.58	1
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	75075.31	2	1	0	92888.52	1
9999	10000	15628319	Walker	792	France	Female	28	4	130142.79	1	1	0	38190.78	0

10000 rows x 14 columns

IN(3):

df.head()

OP(3):

```
Out[3]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.88	1
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.58	0
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	113931.57	1
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	93826.63	0
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	79084.10	0

IN(4):

df.shape

OP(4):

```
Out[4]: (10000, 14)
```

# Univariate,Bivariate and MultiVariate Analysis

## Univariate Analysis

IN[9]:

```
df_france=df.loc[df['Geography']=='France']
```

```
df_spain=df.loc[df['Geography']=='Spain']
```

```
df_germany=df.loc[df['Geography']=='Germany']
```

In [17]:


```
plt.plot(df_france['Balance'],np.zeros_like(df_france['Balance']),'o')
```

```
plt.plot(df_spain['Balance'],np.zeros_like(df_spain['Balance']),'o')
```

```
plt.plot(df_germany['Balance'],np.zeros_like(df_germany['Balance']),'o')
```

```
plt.xlabel('Age')
```


```
plt.show()
```



## Bivariate Analysis

In [18]:

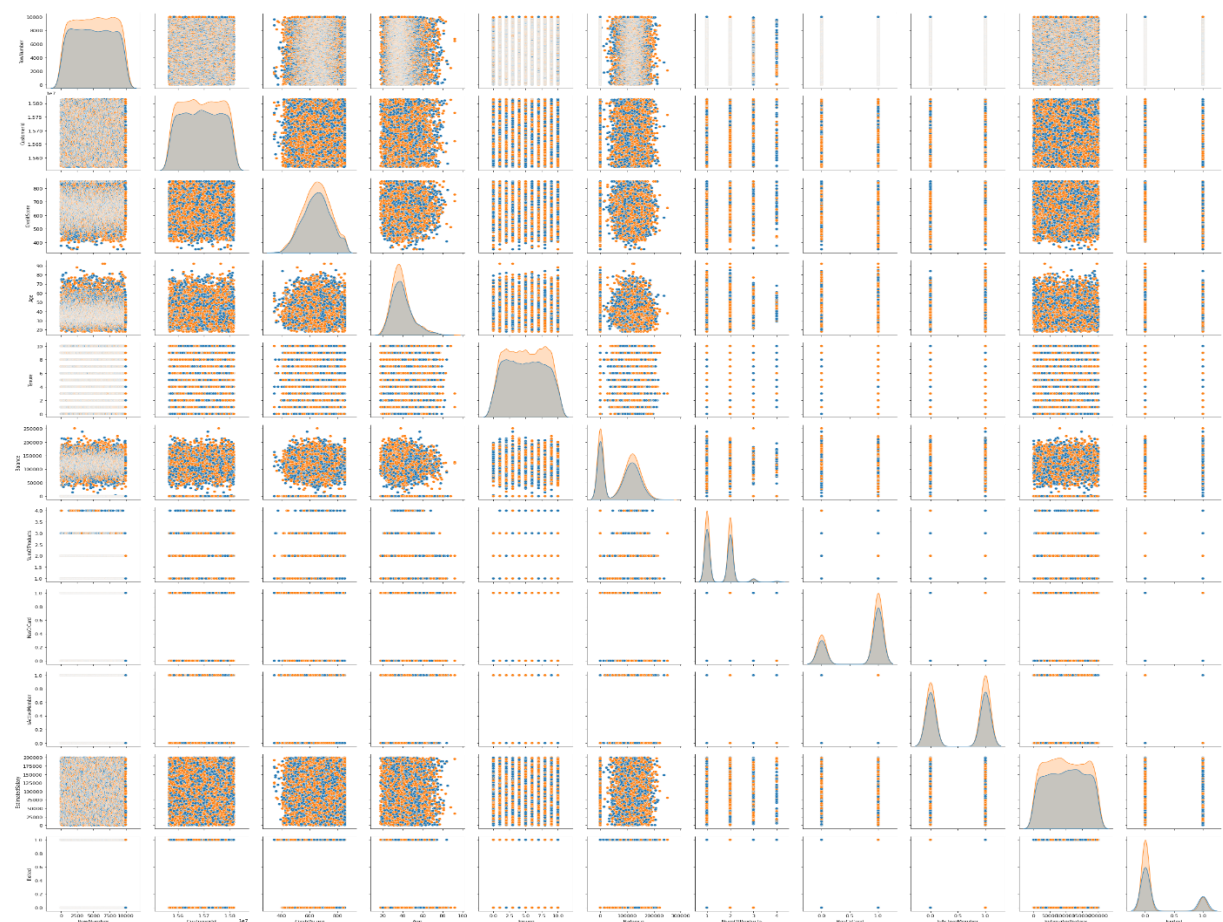
```
sns.FacetGrid(df,hue="Geography",size=5).map(plt.scatter,"Age","Balance").add_legend();  
plt.show()
```



## Multivariate Analysis

In [24]:

```
sns.pairplot(df,hue="Gender",size=3)
```



# Descriptive Statistics

In [29]:  
df.head()

```
Out[29]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.88	1
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.58	0
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	113931.57	1
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	93826.63	0
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	79084.10	0

In [30]:

df.mean() # Get the mean of each column

```
Out[30]:
```

RowNumber	5.000500e+03
CustomerId	1.569094e+07
CreditScore	6.505288e+02
Age	3.892180e+01
Tenure	5.012800e+00
Balance	7.648589e+04
NumOfProducts	1.530200e+00
HasCrCard	7.055000e-01
IsActiveMember	5.151000e-01
EstimatedSalary	1.000902e+05
Exited	2.037000e-01

dtype: float64

In [31]:

df.mean(axis=1) # Get the mean of each row

```
Out[31]:
```

0	1.430602e+06
1	1.440392e+06
2	1.444860e+06
3	1.435993e+06
4	1.449399e+06
...	
9995	1.428483e+06
9996	1.430866e+06
9997	1.421579e+06
9998	1.441922e+06
9999	1.437044e+06

Length: 10000, dtype: float64

In [32]:

df.median() # Get the median of each column

```
Out[32]:
```

RowNumber	5.000500e+03
CustomerId	1.569074e+07
CreditScore	6.520000e+02
Age	3.700000e+01
Tenure	5.000000e+00
Balance	9.719854e+04
NumOfProducts	1.000000e+00
HasCrCard	1.000000e+00
IsActiveMember	1.000000e+00
EstimatedSalary	1.001939e+05
Exited	0.000000e+00

dtype: float64


In [39]:

```
norm_data = pd.DataFrame(np.random.normal(size=100000))
```

```
norm_data.plot(kind="density",  
               figsize=(10,10));
```

```
plt.vlines(norm_data.mean(), # Plot black line at mean  
           ymin=0,  
           ymax=0.4,  
           linewidth=5.0);
```

```
plt.vlines(norm_data.median(), # Plot red line at median  
           ymin=0,  
           ymax=0.4,  
           linewidth=2.0,  
           color="red");
```




In [36]:

```
skewed_data = pd.DataFrame(np.random.exponential(size=100000))
```

```
skewed_data.plot(kind="density",  
                 figsize=(10,10),  
                 xlim=(-1,5));
```

```
plt.vlines(skewed_data.mean(), # Plot black line at mean  
          ymin=0,  
          ymax=0.8,  
          linewidth=5.0);
```

```
plt.vlines(skewed_data.median(), # Plot red line at median  
          ymin=0,  
          ymax=0.8,  
          linewidth=2.0,  
          color="red");
```




In [40]:

```
norm_data = np.random.normal(size=50)
outliers = np.random.normal(15, size=3)
combined_data = pd.DataFrame(np.concatenate((norm_data, outliers), axis=0))

combined_data.plot(kind="density",
                    figsize=(10,10),
                    xlim=(-5,20));

plt.vlines(combined_data.mean(), # Plot black line at mean
            ymin=0,
            ymax=0.2,
            linewidth=5.0);
```

```
plt.vlines(combined_data.median(), # Plot red line at median  
           ymin=0,  
           ymax=0.2,  
           linewidth=2.0,  
           color="red");
```



```
In [42]:  
df.mode()
```



Out[42]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15565701	Smith	850.0	France	Male	37.0	2.0	0.0	1.0	1.0	1.0	24924.92	0.0
1	2	15565706	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2	3	15565714	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
3	4	15565779	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
4	5	15565796	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
9995	9996	15815628	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
9996	9997	15815645	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
9997	9998	15815656	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
9998	9999	15815660	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
9999	10000	15815690	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN

10000 rows × 14 columns

# Measures of Spread

In [43]:

```
max(df["Age"]) - min(df["Age"])
```

Out[43]: 74

In [45]:

```
five_num = [df["Age"].quantile(0),
             df["Age"].quantile(0.25),
             df["Age"].quantile(0.50),
             df["Age"].quantile(0.75),
             df["Age"].quantile(1)]
```

five\_num

Out[45]: [18.0, 32.0, 37.0, 44.0, 92.0]

In [46]:

```
df["Age"].describe()
```

Out[46]:

count	10000.000000
mean	38.921800
std	10.487806
min	18.000000
25%	32.000000
50%	37.000000
75%	44.000000
max	92.000000

Name: Age, dtype: float64

In [47]:

```
df["Age"].quantile(0.75) - df["Age"].quantile(0.25)
```

Out[47]: 12.0

In [49]:

```
df.boxplot(column="Age",  
            return_type='axes',  
            figsize=(8,8))
```

```
plt.text(x=0.74, y=22.25, s="3rd Quartile")
```


```
plt.text(x=0.8, y=18.75, s="Median")
```

```
plt.text(x=0.75, y=15.5, s="1st Quartile")
```

```
plt.text(x=0.9, y=10, s="Min")
```

```
plt.text(x=0.9, y=33.5, s="Max")
```

```
plt.text(x=0.7, y=19.5, s="IQR", rotation=90, size=25);
```



In [50]:

```
df["Age"].var()
```

Out[50]: 109.99408416841683

In [51]:

```
df["Age"].std()
```

```
Out[51]: 10.487806451704609
```

```
In [52]:
```

```
abs_median_devs = abs(df["Age"] - df["Age"].median())
```

```
abs_median_devs.median() * 1.4826
```

```
Out[52]: 8.8956
```

## Skewness and Kurtosis

```
In [53]:
```

```
df["Age"].skew() # Check skewness
```

```
Out[53]: 1.0113202630234552
```

```
In [54]:
```

```
df["Age"].kurt() # Check kurtosis
```

```
Out[54]: 1.3953470615086956
```

```
In [55]:
```

```
norm_data = np.random.normal(size=100000)
```

```
skewed_data = np.concatenate((np.random.normal(size=35000)+2,  
                               np.random.exponential(size=65000)),  
                               axis=0)
```


```
uniform_data = np.random.uniform(0,2, size=100000)
```

```
peaked_data = np.concatenate((np.random.exponential(size=50000),  
                               np.random.exponential(size=50000)*(-1)),  
                               axis=0)
```

```
data_df = pd.DataFrame({"norm":norm_data,  
                        "skewed":skewed_data,  
                        "uniform":uniform_data,  
                        "peaked":peaked_data})
```

```
In [56]:
```

```
data_df.plot(kind="density",  
             figsize=(10,10),  
             xlim=(-5,5));
```



In [57]:

```
data_df.skew()
```

```
Out[57]: norm      -0.007037  
         skewed     1.002549  
         uniform    -0.004434  
         peaked      0.018058  
         dtype: float64
```

In [58]:

```
data_df.kurt()
```

```
Out[58]: norm      -0.009914
         skewed    1.314497
         uniform   -1.201740
         peaked    2.971592
         dtype: float64
```

## Handle the Missing values

In [83]:

```
df=pd.read_csv('/content/Churn_Modelling.csv')
```

```
Out[84]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.88	1
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.58	0
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	113931.57	1
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	93826.63	0
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	79084.10	0

In [84]:

```
df.head()
```

```
Out[86]:
```


	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	False	False	False	False	False	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False	False	False	False	False	False
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
9995	False	False	False	False	False	False	False	False	False	False	False	False	False	False
9996	False	False	False	False	False	False	False	False	False	False	False	False	False	False
9997	False	False	False	False	False	False	False	False	False	False	False	False	False	False
9998	False	False	False	False	False	False	False	False	False	False	False	False	False	False
9999	False	False	False	False	False	False	False	False	False	False	False	False	False	False

10000 rows × 14 columns

In [86]:

```
df.isnull()
```

```
Out[89]: <matplotlib.axes._subplots.AxesSubplot at 0x7f9a987d8290>
```

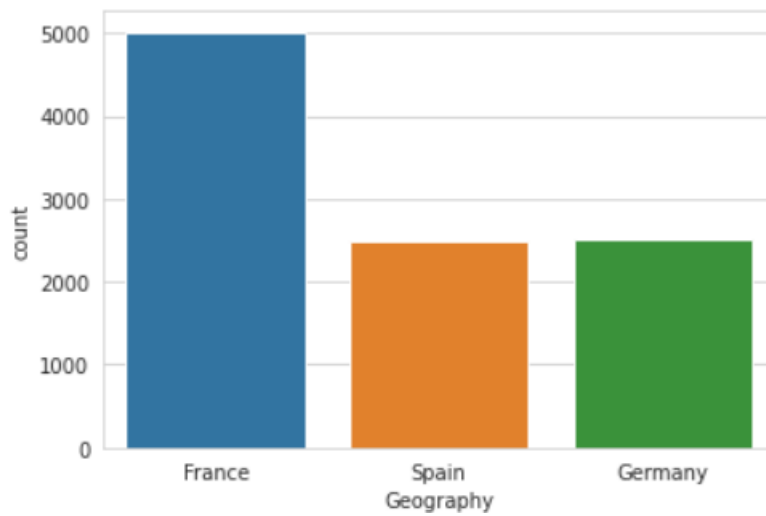


In [93]:

```
sns.set_style('whitegrid')
```

```
sns.countplot(x='Geography',data=df)
```

```
Out[93]: <matplotlib.axes._subplots.AxesSubplot at 0x7f9a92a88850>
```




In [94]:

```
sns.set_style('whitegrid')
```

```
sns.countplot(x='Geography',hue='Gender',data=df,palette='RdBu_r')
```

```
Out[94]: <matplotlib.axes._subplots.AxesSubplot at 0x7f9a92ec10d0>
```




In [96]:

```
sns.set_style('whitegrid')
```

```
sns.countplot(x='Geography',hue='Gender',data=df,palette='rainbow')
```


Out[96]: `<matplotlib.axes._subplots.AxesSubplot at 0x7f9a92afac50>`



In [97]:

```
sns.distplot(df['Age'].dropna(),kde=False,color='darkred',bins=40)
```


Out[97]: `<matplotlib.axes._subplots.AxesSubplot at 0x7f9a98787590>`



In [98]:

```
df['Age'].hist(bins=30,color='darkred',alpha=0.3)
```


Out[98]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7f9a92d64c10>



In [100]:

```
sns.countplot(x='NumOfProducts',data=df)
```


Out[100... <matplotlib.axes.\_subplots.AxesSubplot at 0x7f9a9306f790>



In [101]:

```
df['Age'].hist(color='green',bins=40,figsize=(8,4))
```

Out[101... <matplotlib.axes.\_subplots.AxesSubplot at 0x7f9a90f52d90>





## Cufflinks for plots

In [102]:

```
import cufflinks as cf
cf.go_offline()
```

In []:

```
df['Age'].iplot(kind='hist',bins=30,color='green')
```


## Data Cleaning

In [107]:

```
plt.figure(figsize=(12, 7))
sns.boxplot(x='Gender',y='Age',data=df,palette='winter')
```

Out[107]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f9a90f59450>



In [307]:

```
def impute_age(cols):
    Age = cols[0]
    Pclass = cols[1]

    if pd.isnull(Age):

        if Pclass == 1:
            return 37

        elif Pclass == 2:
            return 29

        else:
            return 24


    else:
        return Age
```

In [122]:

```
sns.heatmap(df.isnull(),yticklabels=False,cbar=False,cmap='viridis')
```

Out[122]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f9a8aa699d0>



In [112]:

```
df.drop('Gender',axis=1,inplace=True)
```

In [114]:

```
df.head()
```

Out[114]:

```
Out[114]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	Hargrave	619	France	42	2	0.00	1	1	1	101348.88	1
1	2	15647311	Hill	608	Spain	41	1	83807.86	1	0	1	112542.58	0
2	3	15619304	Onio	502	France	42	8	159660.80	3	1	0	113931.57	1
3	4	15701354	Boni	699	France	39	1	0.00	2	0	0	93826.63	0
4	5	15737888	Mitchell	850	Spain	43	2	125510.82	1	1	1	79084.10	0

## Converting Categorical Features

In [116]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 10000 entries, 0 to 9999
```

```
Data columns (total 13 columns):
```

#	Column	Non-Null	Count	Dtype
0	RowNumber	10000	non-null	int64
1	CustomerId	10000	non-null	int64
2	Surname	10000	non-null	object
3	CreditScore	10000	non-null	int64
4	Geography	10000	non-null	object
5	Age	10000	non-null	int64
6	Tenure	10000	non-null	int64
7	Balance	10000	non-null	float64
8	NumOfProducts	10000	non-null	int64
9	HasCrCard	10000	non-null	int64
10	IsActiveMember	10000	non-null	int64
11	EstimatedSalary	10000	non-null	float64
12	Exited	10000	non-null	int64

```
dtypes: float64(2), int64(9), object(2)
```

```
memory usage: 1015.8+ KB
```

In [118]:

```
pd.get_dummies(df['Geography'],drop_first=True).head()
```

Out[118]:

	Germany	Spain
0	0	0
1	0	1
2	0	0
3	0	0
4	0	1

In [124]:

```
df.info
```

Out[124]:

```
<bound method DataFrame.info of
ditScore Geography Age Tenure \
0 1 15634602 Hargrave 619 France 42 2
1 2 15647311 Hill 608 Spain 41 1
2 3 15619304 Onio 502 France 42 8
3 4 15701354 Boni 699 France 39 1
4 5 15737888 Mitchell 850 Spain 43 2
... ... ... ... ...
9995 9996 15606229 Obijiaku 771 France 39 5
9996 9997 15569892 Johnstone 516 France 35 10
9997 9998 15584532 Liu 709 France 36 7
9998 9999 15682355 Sabbatini 772 Germany 42 3
9999 10000 15628319 Walker 792 France 28 4
```

	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	0.00	1	1	1	101348.88
1	83807.86	1	0	1	112542.58
2	159660.80	3	1	0	113931.57
3	0.00	2	0	0	93826.63
4	125510.82	1	1	1	79084.10
...	...	...	...	...	...
9995	0.00	2	1	0	96270.64
9996	57369.61	1	1	1	101699.77
9997	0.00	1	0	1	42085.58
9998	75075.31	2	1	0	92888.52
9999	130142.79	1	1	0	38190.78

	Exited
0	1
1	0
2	1
3	0
4	0
...	...
9995	0
9996	0
9997	1
9998	1
9999	0

[10000 rows x 13 columns]>

In [125]:

```
sex = pd.get_dummies(df['Age'],drop_first=True)
embark = pd.get_dummies(df['Balance'],drop_first=True)
```

In [127]:

```
df.drop(['Age', 'HasCrCard', 'Surname', 'CustomerId'],axis=1,inplace=True)
```

In [129]:

```
df.head()
```

Out[129]:

Out[129...]	RowNumber	CreditScore	Geography	Tenure	Balance	NumOfProducts	IsActiveMember	EstimatedSalary	Exited
0	1	619	France	2	0.00	1	1	101348.88	1
1	2	608	Spain	1	83807.86	1	1	112542.58	0
2	3	502	France	8	159660.80	3	0	113931.57	1
3	4	699	France	1	0.00	2	0	93826.63	0
4	5	850	Spain	2	125510.82	1	1	79084.10	0

In [130]:

```
train = pd.concat([df,sex,embark],axis=1)
```

In [131]:

```
train.head()
```

Out[131]:

Out[131...]	RowNumber	CreditScore	Geography	Tenure	Balance	NumOfProducts	IsActiveMember	EstimatedSalary	Exited	19	...	212692.97	212696.32	212778.2	213146.2	214
0	1	619	France	2	0.00	1	1	101348.88	1	0	...	0	0	0	0	
1	2	608	Spain	1	83807.86	1	1	112542.58	0	0	...	0	0	0	0	
2	3	502	France	8	159660.80	3	0	113931.57	1	0	...	0	0	0	0	
3	4	699	France	1	0.00	2	0	93826.63	0	0	...	0	0	0	0	
4	5	850	Spain	2	125510.82	1	1	79084.10	0	0	...	0	0	0	0	

5 rows x 6459 columns



# Find the outliers and replace the outliers

In [147]:

```
dataset= [11,10,12,14,12,15,14,13,15,102,12,14,17,19,107,  
10,13,12,14,12,108,12,11,14,13,15,10,15,12,10,14,13,15,10]
```

## Detecting outlier using Z score

### Using Z score

In [148]:

```
outliers=[]
```

```
def detect_outliers(data):
```

```
    threshold=3
```

```
    mean = np.mean(data)
```

```
    std =np.std(data)
```

```
    for i in data:
```

```
        z_score= (i - mean)/std
```

```
        if np.abs(z_score) > threshold:
```

```
            outliers.append(y)
```

```
    return outliers
```

In [151]:

```
outlier_pt=detect_outliers(dataset)
```

In [152]:

```
outlier_pt
```

Out[152]:

```
Out[152]: [0      101348.88  
1      112542.58  
2      113931.57  
3       93826.63  
4       79084.10  
...  
9995    96270.64  
9996    101699.77  
9997    42085.58  
9998    92888.52  
9999    38190.78  
Name: EstimatedSalary, Length: 10000, dtype: float64, 0      101348.88  
1      112542.58  
2      113931.57  
3       93826.63  
4       79084.10  
...  
9995    96270.64  
9996    101699.77  
9997    42085.58  
9998    92888.52  
9999    38190.78  
Name: EstimatedSalary, Length: 10000, dtype: float64, 0      101348.88  
1      112542.58  
2      113931.57  
3       93826.63  
4       79084.10  
...  
9995    96270.64  
9996    101699.77  
9997    42085.58  
9998    92888.52  
9999    38190.78  
Name: EstimatedSalary, Length: 10000, dtype: float64]
```

In [153]:

```
## Perform all the steps of IQR  
sorted(dataset)
```

Out[153]:

```
Out[153]: [10,  
           10,  
           10,  
           10,  
           10,  
           11,  
           11,  
           12,  
           12,  
           12,  
           12,  
           12,  
           12,  
           12,  
           12,  
           13,  
           13,  
           13,  
           13,  
           14,  
           14,  
           14,  
           14,  
           14,  
           14,  
           15,  
           15,  
           15,  
           15,  
           15,  
           17,  
           19,  
           102,  
           107,  
           108]
```

In [155]:

```
quantile1, quantile3= np.percentile(dataset, [25,75])
```

In [156]:

```
print(quantile1,quantile3)
```

```
12.0 15.0
```

In [157]:

```
## Find the IQR
```

```
iqr_value=quantile3-quantile1
```

```
print(iqr_value)
```

```
3.0
```

In [159]:

```
## Find the lower bound value and the higher bound value
```

```
lower_bound_val = quantile1 -(1.5 * iqr_value)
```

```
upper_bound_val = quantile3 +(1.5 * iqr_value)
```

```
In [160]:
print(lower_bound_val,upper_bound_val)

7.5 19.5
```

## Check for Categorical columns and perform encoding

```
In [161]:
df=pd.read_csv('/content/Churn_Modelling.csv')
```

```
In [162]:
df.head()
```

```
Out[162]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.88	1
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.58	0
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	113931.57	1
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	93826.63	0
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	79084.10	0

```
In [163]:
df_numeric = df[['RowNumber', 'CustomerId', 'CreditScore', 'Age', 'Tenure',
'Balance',
'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary', 'Exited']]
df_categorical = df[['Surname', 'Geography', 'Gender']]
```

```
In [164]:
df_numeric.head()
```

```
Out[164]:
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	619	42	2	0.00	1	1	1	101348.88	1
1	2	15647311	608	41	1	83807.86	1	0	1	112542.58	0
2	3	15619304	502	42	8	159660.80	3	1	0	113931.57	1
3	4	15701354	699	39	1	0.00	2	0	0	93826.63	0
4	5	15737888	850	43	2	125510.82	1	1	1	79084.10	0

```
In [165]:
df_categorical.head()
```

```
Out[165]:
```

	Surname	Geography	Gender
0	Hargrave	France	Female
1	Hill	Spain	Female
2	Onio	France	Female
3	Boni	France	Female
4	Mitchell	Spain	Female

```
In [166]:
print(df['Surname'].unique())
print(df['Geography'].unique())
print(df['Gender'].unique())

['Hargrave' 'Hill' 'Onio' ... 'Kashiwagi' 'Aldridge' 'Burbidge']
['France' 'Spain' 'Germany']
['Female' 'Male']
```

```

In [167]:
from sklearn.preprocessing import LabelEncoder

marry_encoder = LabelEncoder()

In [168]:
marry_encoder.fit(df_categorical['Gender'])

Out[168]:
LabelEncoder()

In [169]:
marry_values = marry_encoder.transform(df_categorical['Gender'])

In [170]:
print("Before Encoding:", list(df_categorical['Gender'][-10:]))
print("After Encoding:", marry_values[-10:])
print("The inverse from the encoding result:",
marry_encoder.inverse_transform(marry_values[-10:]))

Before Encoding: ['Male', 'Female', 'Male', 'Male', 'Female', 'Male', 'Male', 'Male', 'Female', 'Male']
After Encoding: [1 0 1 1 0 1 1 0 1 0]
The inverse from the encoding result: ['Male' 'Female' 'Male' 'Male' 'Female' 'Male' 'Male' 'Female' 'Male' 'Female']

In [171]:
residence_encoder = LabelEncoder()
residence_values =
residence_encoder.fit_transform(df_categorical['Geography'])

print("Before Encoding:", list(df_categorical['Geography'][:5]))
print("After Encoding:", residence_values[:5])
print("The inverse from the encoding result:",
residence_encoder.inverse_transform(residence_values[:5]))

Before Encoding: ['France', 'Spain', 'France', 'France', 'Spain']
After Encoding: [0 2 0 0 2]
The inverse from the encoding result: ['France' 'Spain' 'France' 'France' 'Spain']

In [172]:
from sklearn.preprocessing import OneHotEncoder

gender_encoder = OneHotEncoder()

In [174]:
from sklearn.preprocessing import OneHotEncoder
import numpy as np

gender_encoder = OneHotEncoder()
gender_resaped = np.array(df_categorical['Gender']).reshape(-1, 1)
gender_values = gender_encoder.fit_transform(gender_resaped)

print(df_categorical['Gender'][:5])
print()
print(gender_values.toarray()[:5])
print()
print(gender_encoder.inverse_transform(gender_values[:5])[:5])

0    Female
1    Female
2    Female
3    Female
4    Female
Name: Gender, dtype: object

```



```
[[1. 0.]
 [1. 0.]
 [1. 0.]
 [1. 0.]
 [1. 0.]]
```

```
[['Female']
 ['Female']
 ['Female']
 ['Female']
 ['Female']]
```

In [175]:

```
smoke_encoder = OneHotEncoder()
smoke_resaped = np.array(df_categorical['Surname']).reshape(-1, 1)
smoke_values = smoke_encoder.fit_transform(smoke_resaped)
```

```
print(df_categorical['Surname'][:5])
print()
print(smoke_values.toarray()[:5])
print()
print(smoke_encoder.inverse_transform(smoke_values)[:5])
```

```
0    Hargrave
1      Hill
2      Onio
3      Boni
4  Mitchell
Name: Surname, dtype: object
```

```
[[0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]]
```

```
[['Hargrave']
 ['Hill']
 ['Onio']
 ['Boni']
 ['Mitchell']]
```

In [176]:

```
work_encoder = OneHotEncoder()
work_resaped = np.array(df_categorical['Geography']).reshape(-1, 1)
work_values = work_encoder.fit_transform(work_resaped)
```

```
print(df_categorical['Geography'][:5])
print()
print(work_values.toarray()[:5])
print()
print(work_encoder.inverse_transform(work_values)[:5])
```

```
0    France
1     Spain
2    France
3    France
4     Spain
Name: Geography, dtype: object
```

```
[[1. 0. 0.]
 [0. 0. 1.]
 [1. 0. 0.]
 [1. 0. 0.]]
```

```
[0. 0. 1.]]
```

```
[['France']  
 ['Spain']  
 ['France']  
 ['France']  
 ['Spain']]
```

In [178]:

```
df_categorical_encoded = pd.get_dummies(df_categorical, drop_first=True)  
df_categorical_encoded.head()
```

Out[178]:

	Surname_Abbie	Surname_Abbott	Surname_Abdullah	Surname_Abdulov	Surname_Abel	Surname_Abernathy	Surname_Abramov	Surname_Abramova	Surname_Abramov
0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0	0

5 rows × 2934 columns

In [179]:

```
df_new = pd.concat([df_numeric, df_categorical_encoded], axis=1)  
df_new.head()
```

Out[179]:

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	...	Surname_Zotova	Surname_Zox	Surnam
0	1	15634602	619	42	2	0.00	1	1	1	101348.88	...	0	0	
1	2	15647311	608	41	1	83807.86	1	0	1	112542.58	...	0	0	
2	3	15619304	502	42	8	159660.80	3	1	0	113991.57	...	0	0	
3	4	15701354	699	39	1	0.00	2	0	0	93826.63	...	0	0	
4	5	15737888	850	43	2	125510.82	1	1	1	79084.10	...	0	0	

5 rows × 2945 columns

## Split the data into dependent and independent variables.

In [180]:

```
df=pd.read_csv('/content/Churn_Modelling.csv')
```

In [182]:

```
print(df["Balance"].min())  
print(df["Balance"].max())  
print(df["Balance"].mean())
```

```
0.0
```

```
250898.09
```

```
76485.889288
```

In [183]:

```
print(df.count(0))
```

```
RowNumber      10000
```

```
CustomerId      10000
```

```
Surname         10000
```

```
CreditScore     10000
```

```
Geography       10000
```

```
Gender          10000
```

```
Age            10000
```

```
Tenure         10000
```

```

Balance            10000
NumOfProducts      10000
HasCrCard           10000
IsActiveMember      10000
EstimatedSalary     10000
Exited              10000
dtype: int64
In [184]:
print(df.shape)
(10000, 14)
In [185]:
print(df.size)
140000
In [187]:
X = df.iloc[:, :-1].values
print(X)
[[1 15634602 'Hargrave' ... 1 1 101348.88]
 [2 15647311 'Hill' ... 0 1 112542.58]
 [3 15619304 'Onio' ... 1 0 113931.57]
 ...
 [9998 15584532 'Liu' ... 0 1 42085.58]
 [9999 15682355 'Sabbatini' ... 1 0 92888.52]
 [10000 15628319 'Walker' ... 1 0 38190.78]]
In [271]:
Y = df.iloc[:, -1].values
print(Y)
[1 0 1 ... 1 1 0]

```

## Scale the independent variables

```

In [215]:
df = pd.read_csv('/content/Churn_Modelling.csv')



x = df[['Age', 'Tenure']].values
y = df['Gender'].values

fig, ax = plt.subplots(ncols=2, figsize=(12, 4))

ax[0].scatter(x[:,0], y)
ax[1].scatter(x[:,1], y)

plt.show()

```




In [216]:

```
fig, ax = plt.subplots(figsize=(12, 4))

ax.scatter(x[:,0], y)
ax.scatter(x[:,1], y)
```

Out[216]:

<matplotlib.collections.PathCollection at 0x7f9a8a854ad0>




In [217]:

```
fig, ax = plt.subplots(figsize=(12, 4))

ax.hist(x[:,0])
ax.hist(x[:,1])
```

Out[217]:

```
(array([ 413., 1035., 1048., 1009.,  989., 1012.,  967., 1028., 1025.,
        1474.]),
 array([ 0.,  1.,  2.,  3.,  4.,  5.,  6.,  7.,  8.,  9., 10.]),
 <a list of 10 Patch objects>)
```



In [220]:

```
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
fig, ax = plt.subplots(figsize=(12, 4))


scaler = StandardScaler()
x_std = scaler.fit_transform(x)
```

```
ax.hist(x_std[:,0])
ax.hist(x_std[:,1])
```

Out[220]:

```
(array([ 413., 1035., 1048., 1009., 2001.,    0., 1995.,    0., 1025.,
        1474.]),
 array([-1.73331549, -1.38753759, -1.04175968, -0.69598177, -0.35020386,
        -0.00442596,  0.34135195,  0.68712986,  1.03290776,  1.37868567,
```

```
1.72446358]],
<a list of 10 Patch objects>)
```



In [219]:

```
fig, ax = plt.subplots(figsize=(12, 4))
```

```
scaler = StandardScaler()
```


```
x_std = scaler.fit_transform(x)
```

```
ax.scatter(x_std[:,0], y)
```

```
ax.scatter(x_std[:,1], y)
```

Out[219]:

```
<matplotlib.collections.PathCollection at 0x7f9a8a2fde50>
```



In [221]:

```
fig, ax = plt.subplots(figsize=(12, 4))
```

```
scaler = MinMaxScaler()
```


```
x_minmax = scaler.fit_transform(x)
```

```
ax.hist(x_minmax[:,0])
```

```
ax.hist(x_minmax[:,1])
```

Out[221]:

```
(array([ 413., 1035., 1048., 1009.,  989., 1012.,  967., 1028., 1025.,
        1474.]),
 array([0. , 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1. ]),
<a list of 10 Patch objects>)
```



In [222]:

```
fig, ax = plt.subplots(figsize=(12, 4))
```

```
scaler = MinMaxScaler()
```


```
x_minmax = scaler.fit_transform(x)
```

```
ax.scatter(x_minmax[:,0], y)
```

```
ax.scatter(x_minmax[:,1], y)
```

Out[222]:

<matplotlib.collections.PathCollection at 0x7f9a8a0cae10>



In [223]:

```
fig, ax = plt.subplots(figsize=(12, 4))
```


```
scaler = MinMaxScaler()
```

```
x_minmax = scaler.fit_transform(x)
```

```
ax.scatter(x_minmax[:,0], y)
```

Out[223]:

<matplotlib.collections.PathCollection at 0x7f9a8a0caf10>



In [224]:

```
fig, ax = plt.subplots(figsize=(12, 4))
```

```

scaler = MinMaxScaler()
x_minmax = scaler.fit_transform(x)

ax.hist(x_minmax[:,0])


```

Out[224]:

```

(array([ 611., 2179., 3629., 1871.,  910.,  441.,  208.,  127.,   20.,
         4.]),
 array([0. , 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1. ]),
 <a list of 10 Patch objects>)

```



In [227]:

```

from sklearn.model_selection import train_test_split
from sklearn.pipeline import Pipeline
from sklearn.linear_model import SGDRegressor
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import mean_absolute_error
import sklearn.metrics as metrics

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

# Import Data
df = pd.read_csv('/content/Churn_Modelling.csv')
x = df[['Age', 'Tenure']].values
y = df['Balance'].values

# Split into a training and testing set
X_train, X_test, Y_train, Y_test = train_test_split(x, y)

# Define the pipeline for scaling and model fitting
pipeline = Pipeline([
    ("MinMax Scaling", MinMaxScaler()),
    ("SGD Regression", SGDRegressor())
])

# Scale the data and fit the model
pipeline.fit(X_train, Y_train)

# Evaluate the model
Y_pred = pipeline.predict(X_test)
print('Mean Absolute Error: ', mean_absolute_error(Y_pred, Y_test))
print('Score', pipeline.score(X_test, Y_test))

Mean Absolute Error:  57120.533393590835
Score 0.0004207814312172653

```

# Split the data into training and testing

In [267]:

```
dataset = pd.read_csv('/content/Churn_Modelling.csv')
print(dataset)
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age
\							
0	1	15634602	Hargrave	619	France	Female	42
1	2	15647311	Hill	608	Spain	Female	41
2	3	15619304	Onio	502	France	Female	42
3	4	15701354	Boni	699	France	Female	39
4	5	15737888	Mitchell	850	Spain	Female	43
...	...	...	...	...	...	...	...
9995	9996	15606229	Obijiaku	771	France	Male	39
9996	9997	15569892	Johnstone	516	France	Male	35
9997	9998	15584532	Liu	709	France	Female	36
9998	9999	15682355	Sabbatini	772	Germany	Male	42
9999	10000	15628319	Walker	792	France	Female	28

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	
4	2	125510.82	1	1	1	
...	...	...	...	...	...	...
9995	5	0.00	2	1	0	
9996	10	57369.61	1	1	1	
9997	7	0.00	1	0	1	
9998	3	75075.31	2	1	0	
9999	4	130142.79	1	1	0	

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0
...	...	...
9995	96270.64	0
9996	101699.77	0
9997	42085.58	1
9998	92888.52	1
9999	38190.78	0

[10000 rows x 14 columns]

In [287]:

```
dataset.drop(["HasCrCard"],axis=1,inplace=True)
```

In [288]:

```
print(dataset.shape)#no. of rows and columne
print(dataset.head(10))
```

(10000, 7)

	CustomerId	CreditScore	Age	Tenure	Balance	IsActiveMember	\
0	15634602	619	42	2	0.00	1	
1	15647311	608	41	1	83807.86	1	
2	15619304	502	42	8	159660.80	0	
3	15701354	699	39	1	0.00	0	
4	15737888	850	43	2	125510.82	1	



5	15574012	645	44	8	113755.78	0
6	15592531	822	50	7	0.00	1
7	15656148	376	29	4	115046.74	0
8	15792365	501	44	4	142051.07	1
9	15592389	684	27	2	134603.88	1

```

EstimatedSalary
0      101348.88
1      112542.58
2      113931.57
3       93826.63
4       79084.10
5      149756.71
6       10062.80
7      119346.88
8       74940.50
9       71725.73

```

In [289]:

```

X=dataset.iloc[:, :-1].values
X

```

Out[289]:

```

array([[1.5634602e+07, 6.1900000e+02, 4.2000000e+01, 2.0000000e+00,
        0.0000000e+00, 1.0000000e+00],
       [1.5647311e+07, 6.0800000e+02, 4.1000000e+01, 1.0000000e+00,
        8.3807860e+04, 1.0000000e+00],
       [1.5619304e+07, 5.0200000e+02, 4.2000000e+01, 8.0000000e+00,
        1.5966080e+05, 0.0000000e+00],
       ...,
       [1.5584532e+07, 7.0900000e+02, 3.6000000e+01, 7.0000000e+00,
        0.0000000e+00, 1.0000000e+00],
       [1.5682355e+07, 7.7200000e+02, 4.2000000e+01, 3.0000000e+00,
        7.5075310e+04, 0.0000000e+00],
       [1.5628319e+07, 7.9200000e+02, 2.8000000e+01, 4.0000000e+00,
        1.3014279e+05, 0.0000000e+00]])

```

In [290]:

```

Y=dataset.iloc[:, -1].values
Y

```

Out[290]:

```

array([101348.88, 112542.58, 113931.57, ..., 42085.58, 92888.52,
       38190.78])

```

In [291]:

```

from sklearn.model_selection import train_test_split
X_train,X_test,Y_train,Y_test = train_test_split( X, Y, test_size = 0.25,
random_state = 0 )

```

In [306]:

```

from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
print(X_train)

[[-1.34333028 -0.73550706  0.01526571  0.00886037  0.67316003 -1.03446007]
 [ 1.55832963  1.02442719 -0.65260917  0.00886037 -1.20772417 -1.03446007]
 [-0.65515619  0.80829492 -0.46178778  1.39329338 -0.35693706  0.96668786]
 ...
 [-1.63542994  0.90092304 -0.36637708  0.00886037  1.36657199 -1.03446007]
 [-0.38540456 -0.62229491 -0.08014499  1.39329338 -1.20772417  0.96668786]
 [-1.37829524 -0.28265848  0.87396199 -1.37557264  0.51741687 -1.03446007]]

```

In [305]:

```
print(X_test)
```

```
[[-1.05852196 -0.55025082 -0.36637708  1.04718513  0.88494297  0.96668786]
 [-0.51554728 -1.31185979  0.11067641 -1.02946438  0.43586703 -1.03446007]
 [-0.8058485   0.57157862  0.3014978   1.04718513  0.31486378  0.96668786]
 ...
 [ 0.25326371  1.95070838  0.01526571 -1.37557264  0.30819395 -1.03446007]
 [-0.17836122  0.29369426 -0.08014499  0.70107688  0.55698791 -1.03446007]
 [ 0.40190663  0.870047   -0.74801987 -0.68335613  0.7006957  -1.03446007]]
```